

Piyush Chaudhari

piyushchaudhari151001@gmail.com | 9923019644 | LinkedIn | Github

Experience

Media.net

07/2023 – Present

Data Application Developer 2

- Developed **Spark (PySpark/Scala)** ETL pipelines to migrate existing **Hive**-based datasets to **Apache Iceberg** tables, improving query performance and **reducing data processing time by 20%**.
- Contributed to the design and development of **Spring-based** applications, including an adtech entity management system and a data replication tool, with a focus on scalability and reliability.
- Designed and implemented database pipelines for **Ads.txt validation** along with **Grafana dashboards** for monitoring and alerting, eliminating bids on unvalidated inventory and enabling monetization of previously untargeted but valid ad requests.
- Architected database schemas and implemented stored procedure logic for a **programmatic ad-tech entity system**, utilizing **bidirectional replication** to support high availability and controlled write failover.

Projects

Unified Reporting

- Designed and implemented a **Scala / Apache Spark** batch ETL pipeline that ingests windows from JDBC metadata (rather than fixed dates), and loads unified hourly/daily analytics into a data lake using Apache Iceberg (Spark catalog, REST-backed metadata, custom read/write paths).
- Developed a unified data aggregation pipeline combining multiple metrics and dimensions into standardized reporting datasets for publishers.

Magneto - Real-Time Data Replication System

- Contributed to a Spring Boot **REST API** that monitors **Kafka Connect**-based change data capture (CDC) between MySQL and Microsoft SQL Server, aggregating connector status, Kafka consumer lag, and database-side signals into health and lag responses for operations teams.
- Supported operational automation by helping implement connector lifecycle actions (create, pause, resume, restart, task restart) against Kafka Connect, reducing manual intervention during incident response and connector maintenance.
- Worked on heterogeneous RDBMS replication tooling combining JDBC, stored procedure-driven source lag checks, and **Debezium-oriented** connector configuration patterns (SQL Server and MySQL paths) as part of a cross-functional engineering effort.

AdTech Entity Management System

- Designed and implemented stored procedures for core modules including **Feature Management**, **Ads.txt Crawling**, **A/B Testing**, and **AdTech Entity Management**, supporting client-facing workflows.
- Led end-to-end migration of **100+** stored procedures from SQL Server to MySQL, executing a **phased go-live strategy** to minimize downtime.
- Designed and tuned partitioned ads.txt datasets (**1B+ rows**) used for schain validation, improving query efficiency via partition pruning and stable verification of supply-chain objects.

Skills

- **Concepts:** RESTful APIs, Microservices, OOPS, DBMS, System Design, DSA, Problem Solving
- **Technologies:** Spring, Spring Boot, MySQL, SQL Server, Postgres, Git, Docker, Spark, Airflow
- **Programming Languages:** Java, SQL, Python, C, C++, Scala

Education

Mumbai University, India

08/2019 – 04/2023

Bachelor of Technology, Information Technology **GPA: 9.23/10.0**